

emp

September 28, 2023

#

Employee attrition (2512)

0.0.1 Imports

```
[ ]: import pandas as pd
import seaborn as sns
import numpy as np
import warnings
import matplotlib.pyplot as plt
warnings.filterwarnings('ignore')
from sklearn.preprocessing import StandardScaler
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
from sklearn.tree import DecisionTreeClassifier
```

0.0.2 Data Loading

```
[ ]: df = pd.read_csv("Emp.csv")
```

0.0.3 Null checking

```
[ ]: # df.shape
df.isnull().sum()
```

```
[ ]: Age                0
Attrition              0
BusinessTravel         0
DailyRate              0
Department             0
DistanceFromHome       0
Education              0
EducationField          0
EmployeeCount           0
```

EmployeeNumber	0
EnvironmentSatisfaction	0
Gender	0
HourlyRate	0
JobInvolvement	0
JobLevel	0
JobRole	0
JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0
PercentSalaryHike	0
PerformanceRating	0
RelationshipSatisfaction	0
StandardHours	0
StockOptionLevel	0
TotalWorkingYears	0
TrainingTimesLastYear	0
WorkLifeBalance	0
YearsAtCompany	0
YearsInCurrentRole	0
YearsSinceLastPromotion	0
YearsWithCurrManager	0
dtype: int64	

Clearly, we dont have any null values to be processed in the dataset

Data Visualizations

```
[ ]: # Visualization 1: Bar Chart for Attrition

plt.figure(figsize=(6, 4))
sns.countplot(x='Attrition', data=df)
plt.title('Attrition Distribution')
plt.xlabel('Attrition')
plt.ylabel('Count')
plt.show()

# Visualization 2: Histogram for Age

plt.figure(figsize=(8, 4))
sns.histplot(df['Age'], bins=20, kde=True)
plt.title('Age Distribution')
plt.xlabel('Age')
```

```

plt.ylabel('Count')
plt.show()

# Visualization 3: Count Plot for Business Travel

plt.figure(figsize=(8, 4))
sns.countplot(x='BusinessTravel', data=df)
plt.title('Business Travel Distribution')
plt.xlabel('Business Travel')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()



# Visualization 4: Correlation Heatmap for Numerical Features


numerical_features = df.select_dtypes(include=['int64', 'float64']).columns
correlation_matrix = df[numerical_features].corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap for Numerical Features')
plt.show()

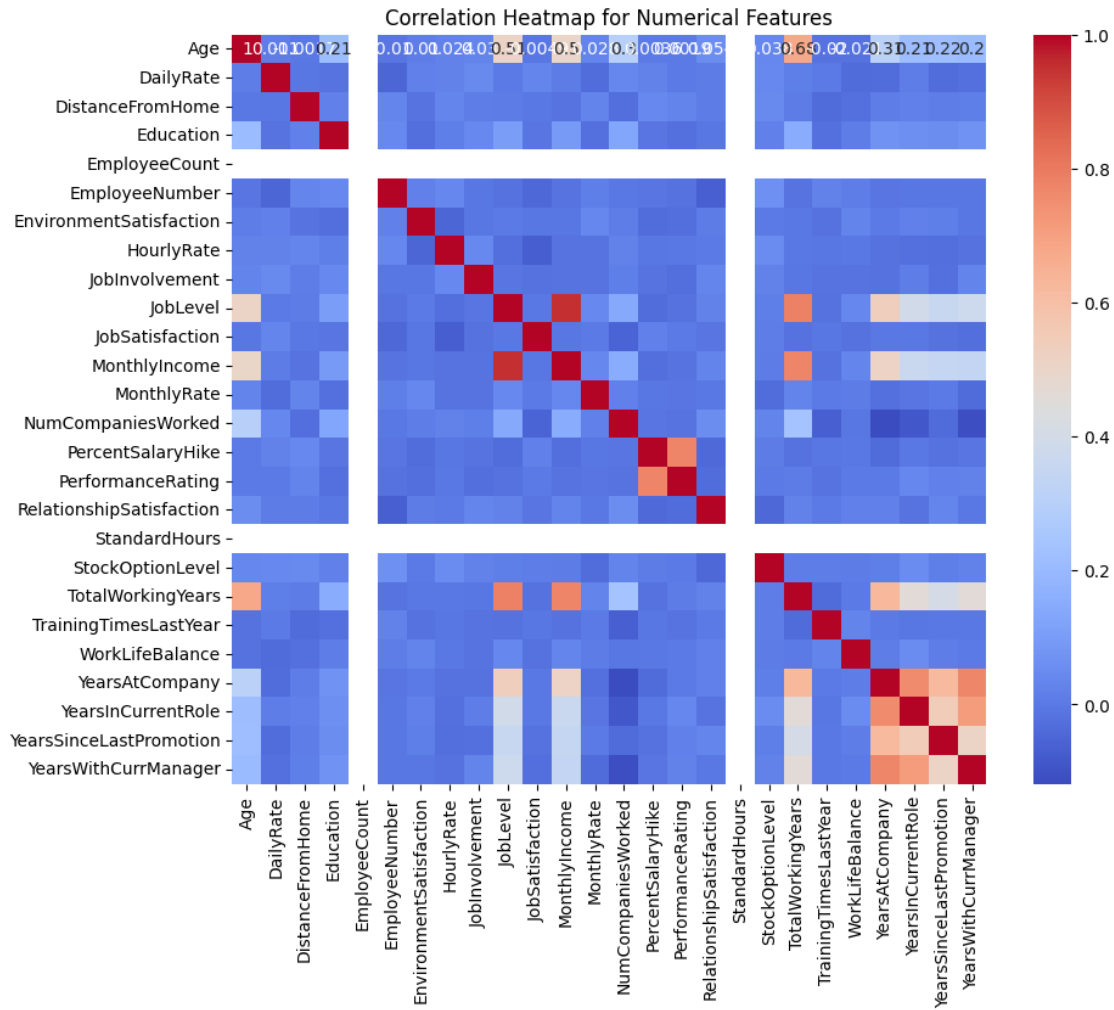
# Visualization 5: Attrition by Department


plt.figure(figsize=(8, 4))
sns.countplot(x='Department', hue='Attrition', data=df)
plt.title('Attrition by Department')
plt.xlabel('Department')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.legend(title='Attrition', loc='upper right')
plt.show()

```









0.0.4 Outlier Detection

```
[ ]: numerical_features = ["Age", "DailyRate", "DistanceFromHome", "HourlyRate",
    ↪ "MonthlyIncome", "MonthlyRate", "NumCompaniesWorked", "PercentSalaryHike",
    ↪ "TotalWorkingYears", "YearsAtCompany", "YearsInCurrentRole",
    ↪ "YearsSinceLastPromotion", "YearsWithCurrManager"]

z_scores = np.abs(stats.zscore(df[numerical_features]))

#Z-score calc,
threshold = 5

# Find and print the indices of outliers
outlier_indices = (z_scores > threshold).any(axis=1)
outliers = df[outlier_indices]

print("Outliers:")
print(outliers)
```

Outliers:

Age Attrition BusinessTravel DailyRate Department \

```

126      58      Yes  Travel_Rarely      147  Research & Development

      DistanceFromHome  Education  EducationField  EmployeeCount  \
126              23              4      Medical              1

      EmployeeNumber  ...  RelationshipSatisfaction  StandardHours  \
126              165  ...              4              80

      StockOptionLevel  TotalWorkingYears  TrainingTimesLastYear  \
126              1              40              3

      WorkLifeBalance  YearsAtCompany  YearsInCurrentRole  \
126              2              40              10

      YearsSinceLastPromotion  YearsWithCurrManager
126              15              6

[1 rows x 35 columns]

```

Hence for the given threshold, we can see that there is only one outlier here, so we shall not worry about it.

0.0.5 Splitting into Dependent and Independent variables

```

[ ]: X = df.drop('Attrition', axis=1)
     y = df['Attrition']
     X

[ ]:
     Age      BusinessTravel  DailyRate      Department  \
0      41      Travel_Rarely      1102      Sales
1      49  Travel_Frequently      279  Research & Development
2      37      Travel_Rarely      1373  Research & Development
3      33  Travel_Frequently      1392  Research & Development
4      27      Travel_Rarely      591  Research & Development
...  ...
1465   36  Travel_Frequently      884  Research & Development
1466   39      Travel_Rarely      613  Research & Development
1467   27      Travel_Rarely      155  Research & Development
1468   49  Travel_Frequently      1023      Sales
1469   34      Travel_Rarely      628  Research & Development

      DistanceFromHome  Education  EducationField  EmployeeCount  \
0              1              2  Life Sciences              1
1              8              1  Life Sciences              1
2              2              2      Other              1
3              3              4  Life Sciences              1
4              2              1      Medical              1

```


...
1465	23	2	Medical	1
1466	6	1	Medical	1
1467	4	3	Life Sciences	1
1468	2	3	Medical	1
1469	8	3	Medical	1

	EmployeeNumber	EnvironmentSatisfaction	...	RelationshipSatisfaction	\
0	1		2	...	1
1	2		3	...	4
2	4		4	...	2
3	5		4	...	3
4	7		1	...	4
...	
1465	2061		3	...	3
1466	2062		4	...	1
1467	2064		2	...	2
1468	2065		4	...	4
1469	2068		2	...	1

	StandardHours	StockOptionLevel	TotalWorkingYears	\
0	80	0	8	
1	80	1	10	
2	80	0	7	
3	80	0	8	
4	80	1	6	
...	
1465	80	1	17	
1466	80	1	9	
1467	80	1	6	
1468	80	0	17	
1469	80	0	6	

	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
0	0	1	6	
1	3	3	10	
2	3	3	0	
3	3	3	8	
4	3	3	2	
...	
1465	3	3	5	
1466	5	3	7	
1467	0	3	6	
1468	3	2	9	
1469	3	4	4	

YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager
--------------------	-------------------------	----------------------

0	4	0	5
1	7	1	7
2	0	0	0
3	7	3	0
4	2	2	2
...
1465	2	0	3
1466	7	1	7
1467	2	0	3
1468	6	0	8
1469	3	1	2

[1470 rows x 34 columns]

0.0.6 Encoding categorical variables

```
[ ]: categorical_columns = ["BusinessTravel", "Department", "EducationField",
    ↪ "Gender", "JobRole", "MaritalStatus", "OverTime", "Over18"]

X_encoded = pd.get_dummies(X, columns=categorical_columns)
```

Feature Scaling is not required here, since only Decision trees and Logistic regression are being done neither of which are scale sensitive

0.0.7 Splitting into testing and training

```
[ ]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.
    ↪ 2, random_state=42)
y_train = y_train.map({'Yes': 1, 'No': 0})
```

```
[ ]: X_train
# y_train.shape
# X_test.shape
# y_test.shape
# y_train
```

```
[ ]:      Age  DailyRate  DistanceFromHome  Education  EmployeeCount  \
1097    24         350             21           2             1
727     18         287              5           2             1
254     29        1247             20           2             1
1175    39         492             12           3             1
1341    31         311             20           3             1
...     ...         ...             ...           ...           ...
1130    35         750             28           3             1
1294    41         447              5           3             1
```

860	22	1256	3	4	1
1459	29	1378	13	2	1
1126	50	264	9	3	1

	EmployeeNumber	EnvironmentSatisfaction	HourlyRate	JobInvolvement	\
1097	1551	3	57	2	
727	1012	2	73	3	
254	349	4	45	3	
1175	1654	4	66	3	
1341	1881	2	89	3	
...	
1130	1596	2	46	4	
1294	1814	2	85	4	
860	1203	3	48	2	
1459	2053	4	46	2	
1126	1591	3	59	3	

	JobLevel	...	JobRole_Manufacturing Director	\
1097	1	...	False	
727	1	...	False	
254	2	...	False	
1175	2	...	True	
1341	2	...	False	
...	
1130	2	...	False	
1294	2	...	False	
860	1	...	False	
1459	2	...	False	
1126	5	...	False	

	JobRole_Research Director	JobRole_Research Scientist	\
1097	False	False	
727	False	True	
254	False	False	
1175	False	False	
1341	False	False	
...	
1130	False	False	
1294	False	False	
860	False	True	
1459	False	False	
1126	False	False	

	JobRole_Sales Executive	JobRole_Sales Representative	\
1097	False	False	
727	False	False	
254	True	False	

1175	False	False
1341	False	False
...
1130	False	False
1294	False	False
860	False	False
1459	False	False
1126	False	False

	MaritalStatus_Divorced	MaritalStatus_Married	MaritalStatus_Single \
1097	True	False	False
727	False	False	True
254	True	False	False
1175	False	True	False
1341	True	False	False
...
1130	False	True	False
1294	False	False	True
860	False	True	False
1459	False	True	False
1126	False	True	False

	OverTime_No	OverTime_Yes
1097	True	False
727	True	False
254	True	False
1175	True	False
1341	True	False
...
1130	True	False
1294	True	False
860	False	True
1459	False	True
1126	False	True

[1176 rows x 55 columns]

```
[ ]: logistic_regression_model = LogisticRegression(random_state=42)

# Fit the model on the training data
logistic_regression_model.fit(X_train, y_train)

# Predict on the test data
y_pred = logistic_regression_model.predict(X_test)
y_pred = np.where(y_pred == 0, 'No', 'Yes')
```

```

# Convert y_test to string labels if it's not already
y_test = y_test.astype(str)

# Now both y_test and y_pred have string labels

# Evaluate the model's performance using y_test
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred)

# Print the results
print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)

```

Accuracy: 0.8673469387755102

Confusion Matrix:

```

[[255  0]
 [ 39  0]]

```

Classification Report:

	precision	recall	f1-score	support
No	0.87	1.00	0.93	255
Yes	0.00	0.00	0.00	39
accuracy			0.87	294
macro avg	0.43	0.50	0.46	294
weighted avg	0.75	0.87	0.81	294

```

[ ]: decision_tree_model = DecisionTreeClassifier(random_state=42)
decision_tree_model.fit(X_train, y_train)

# Predict on the test data
y_true = y_test.map({'No': 0, 'Yes': 1})
y_pred = decision_tree_model.predict(X_test)
accuracy = accuracy_score(y_true, y_pred)
conf_matrix = confusion_matrix(y_true, y_pred)
class_report = classification_report(y_true, y_pred)

# Print the results
print("Accuracy:", accuracy)
print("Confusion Matrix:\n", conf_matrix)
print("Classification Report:\n", class_report)

```

Accuracy: 0.7653061224489796

Confusion Matrix:

[[217 38]

[31 8]]

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.85	0.86	255
1	0.17	0.21	0.19	39
accuracy			0.77	294
macro avg	0.52	0.53	0.53	294
weighted avg	0.78	0.77	0.77	294